

📉💰 A Comparative Study of Ridge and Lasso Regularization Based on MSE Performance on Advertising Data

This example explores how **L1 (Lasso)** and **L2 (Ridge)** regularization techniques help improve prediction accuracy and avoid overfitting when building a regression model to predict **Sales Revenue** from various **advertising spending channels**.

🔍 Scenario: Predicting Sales from Ad Spend

You have the following features:

- **TV Advertising Spend**
- **Radio Advertising Spend**
- **Newspaper Advertising Spend**
- **Social Media Advertising Spend**
- **Influencer Marketing Spend**
- **Search Engine Marketing Spend**

Your linear regression model aims to predict:

Sales Revenue = w1*TV + w2*Radio + w3*Newspaper + w4*Socia Media + w5*Influencer + w6*Search Engine + b

⚠️ Without Regularization (Basic Linear Regression)

A plain linear regression model might **overfit** the data if:

- Some channels don't strongly influence sales
- There's overlap or multicollinearity between ad channels
- Data contains noise or outliers

Overfitting means the model performs well on training data but poorly on new/unseen data.

📊 L2 Regularization (Ridge)

Suppose Ridge regression estimates:

- w1 = 3.2 (TV)
- w2 = 2.0 (Radio)
- w3 = 1.1 (Newspaper)
- w4 = 1.8 (Social Media)
- w5 = 0.9 (Influencer)
- w6 = 2.5 (Search Engine)

Ridge **shrinks** all coefficients but **retains all features**, assuming each has some relevance.

✅ **Best when all ad channels contribute to sales performance.**

✂️ L1 Regularization (Lasso)

Now suppose Lasso regression results in:

- w1 = 3.5 (TV)
- w2 = 1.9 (Radio)
- w3 = 0 (Newspaper)
- w4 = 2.0 (Social Media)
- w5 = 0 (Influencer)
- w6 = 2.7 (Search Engine)

Lasso sets some coefficients to **zero**, removing channels that don't significantly contribute to sales.

✅ **Great when you want simpler models and automatic feature selection.**

💡 Intuition Recap

Regularization	What it does	When to use
Ridge (L2)	Shrinks all coefficients (keeps all features)	When all ad channels matter to some extent
Lasso (L1)	Eliminates unimportant features (zero weights)	When simplifying the model or selecting top ads

With out Regulation

```
import pandas as pd

# Load the dataset
data = pd.read_excel("/content/Sales.xlsx")

data

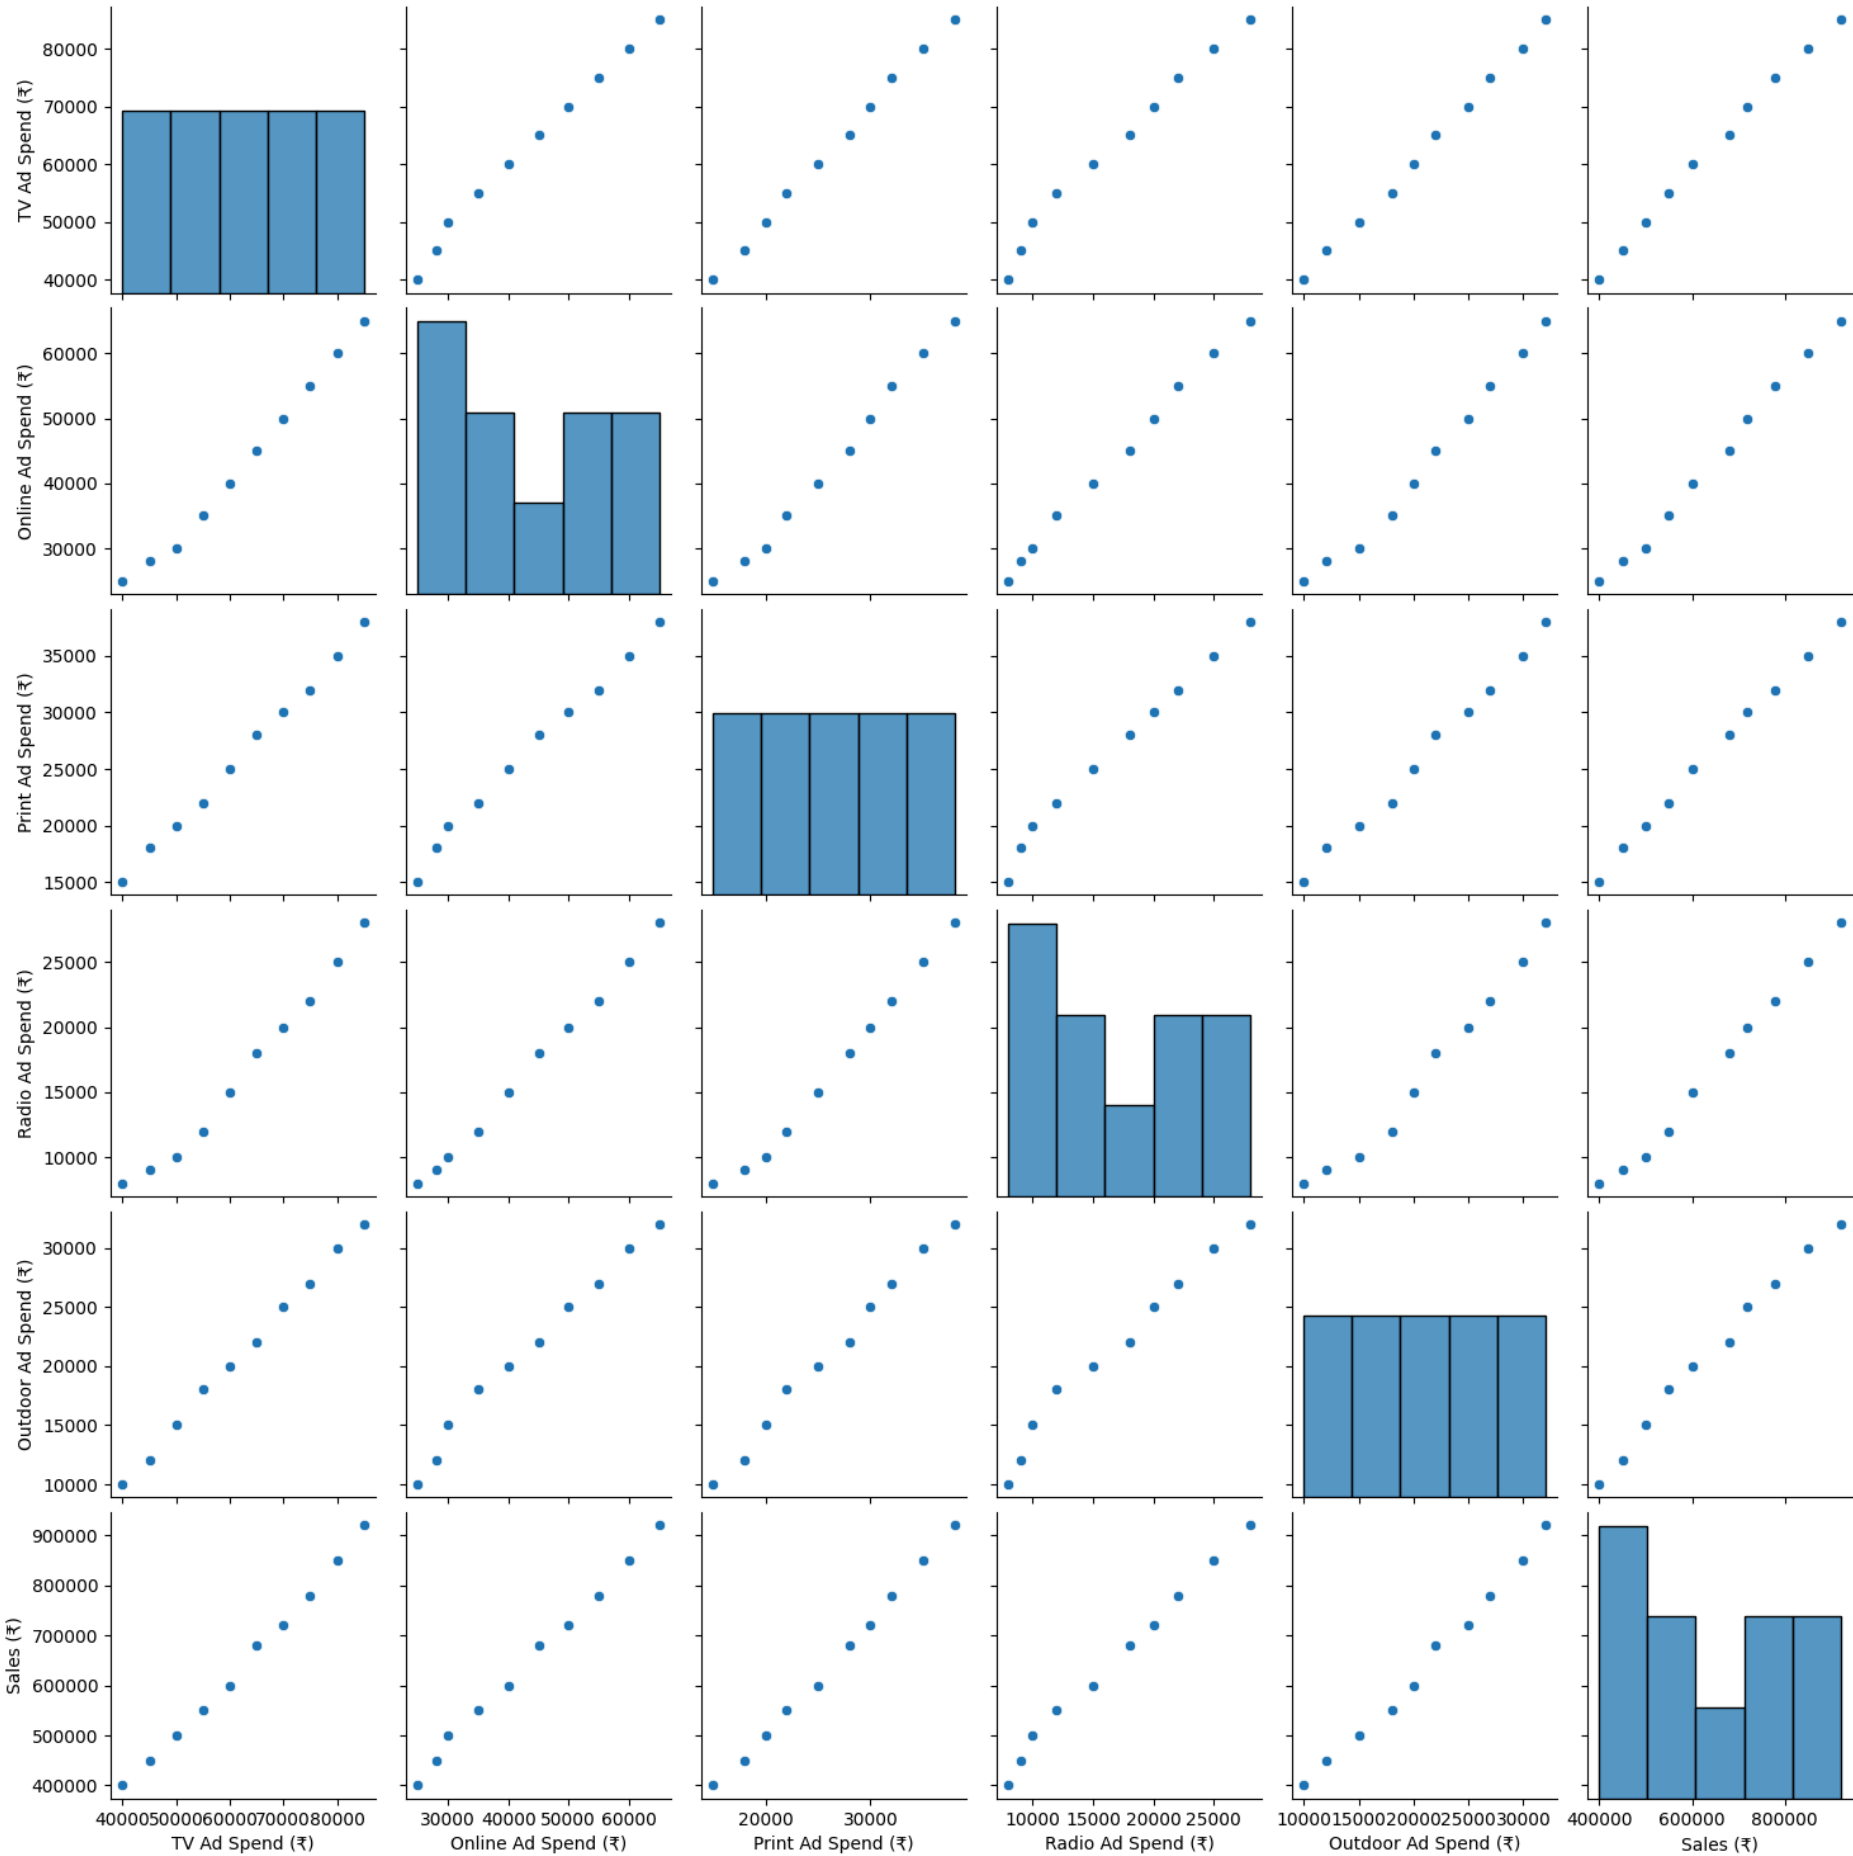
TV Ad Spend (₹)  Online Ad Spend (₹)  Print Ad Spend (₹)  Radio Ad Spend (₹)  Outdoor Ad Spend (₹)  Sales (₹)
0      50000      30000      20000      10000      15000      500000
1      60000      40000      25000      15000      20000      600000
2      70000      50000      30000      20000      25000      720000
3      40000      25000      15000      8000      10000      400000
4      80000      60000      35000      25000      30000      850000
5      55000      35000      22000      12000      18000      550000
6      65000      45000      28000      18000      22000      680000
7      75000      55000      32000      22000      27000      780000
8      45000      28000      18000      9000      12000      450000
9      85000      65000      38000      28000      32000      920000

data.columns

Index(['TV Ad Spend (₹)', 'Online Ad Spend (₹)', 'Print Ad Spend (₹)',
      'Radio Ad Spend (₹)', 'Outdoor Ad Spend (₹)', 'Sales (₹)'],
      dtype='object')

import seaborn as sn
sn.pairplot(data)
```

<seaborn.axisgrid.PairGrid at 0x7aa44b3047d0>



```
# Features and target
X = data[['TV Ad Spend (₹)', 'Online Ad Spend (₹)', 'Print Ad Spend (₹)', 'Radio Ad Spend (₹)', 'Outdoor Ad Spend (₹)']]
y = data["Sales (₹)"]
```

```
# Split the dataset
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Linear Regression (no regularization)
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)
```

LinearRegression()

```
print("Linear Regression Coefficients:", lr.coef_)
```

```
Linear Regression Coefficients: [ 4.82058824 -0.62058824 11.63235294  8.36764706 -5.          ]
```

```
from sklearn.metrics import mean_squared_error
y_pred = lr.predict(X_test)
mse_No_regulation = mean_squared_error(y_test, y_pred)
mse_No_regulation
```

```
85104913.49480934
```

L1 Regulation

```
# Linear Regression (L1 regularization)
from sklearn.linear_model import Lasso
lasso = Lasso(alpha=10000)

lasso.fit(X_train, y_train)
```

